



V3-noSQ-SVDQ decomposition: where the gain comes from (LLM SVDQ vs DiT SVDQ — single-seed, existing data only, noise ± 5 -10pp/suite)



Decomposition findings (Δ vs A0 = Method C, 4-avg):

- L_only (LLM SVDQ only, DiT unchanged) $\rightarrow +3.0\text{pp}$ $\leftarrow$ main contributor (object +11, goal +3)
- D_only (DiT SVDQ only, LLM unchanged) $\rightarrow +0.75\text{pp}$ / -0.75pp $\leftarrow$ within single-seed noise ($\pm 5\text{pp}$)
- LD_both (LLM SVDQ + DiT SVDQ) $\rightarrow +0.25\text{pp}$ $\leftarrow$ no synergy; long regresses (-4pp vs L_only)

Read: V3-noSQ-SVDQ's gain comes mostly from the LLM side; DiT SVDQ is neutral, and stacking both shows mild anti-synergy on long-horizon.